


**SUPER MARIO
BROS**

**INTELLIGENT
AGENT**






**TRAINED BY REINFORCEMENT
LEARNING**

**NOURIN AHMED
KALYAN VENKATESH POLUDASU
SIDDHARTHA PITCHIKA**



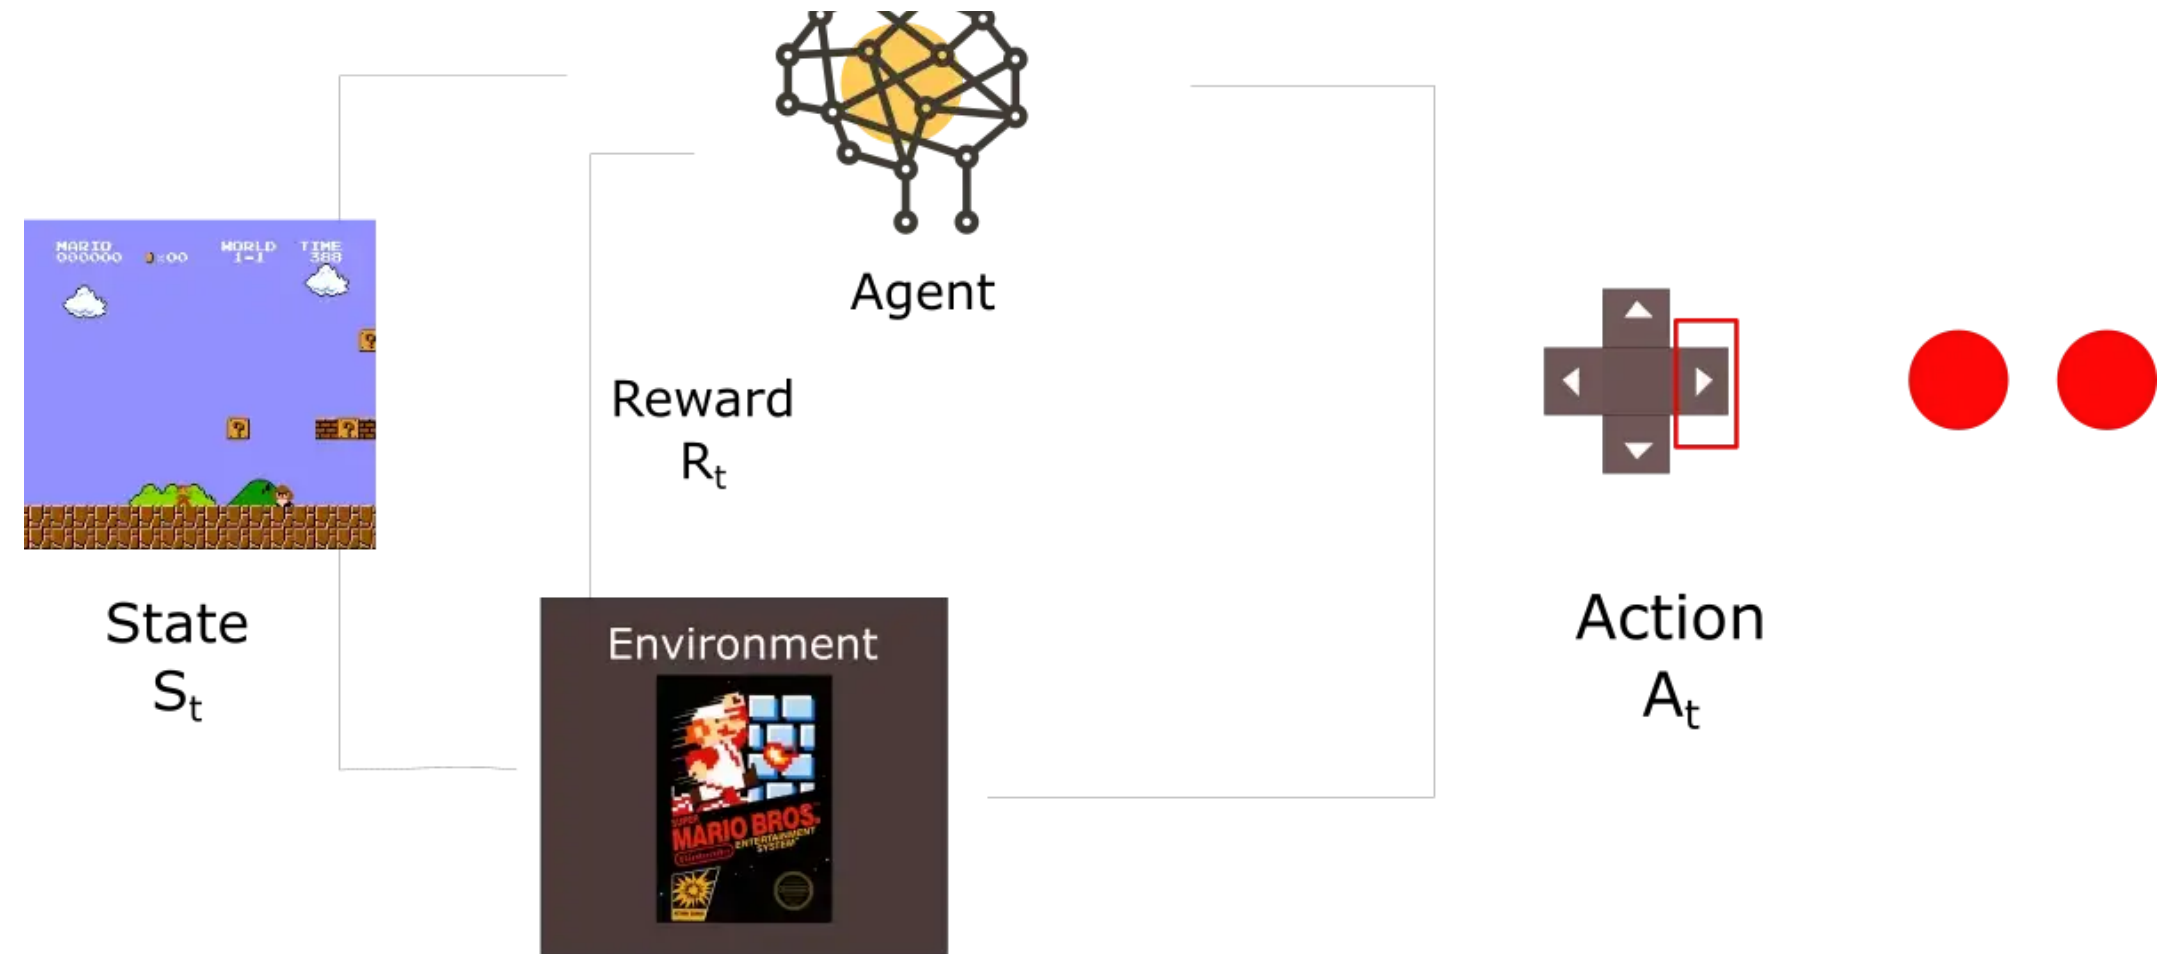
PROBLEM STATEMENT

- Reinforcement Algorithms have been used to solve Atari games.
- There are not much work done on solving NES(Nintendo Platform games) games with RL.
- Our approach is to implement different RL algorithms to train an intelligent Super Mario Agent.



WHAT IS REINFORCEMENT LEARNING?

- Agent
- Environment
- State
- Action
- Reward

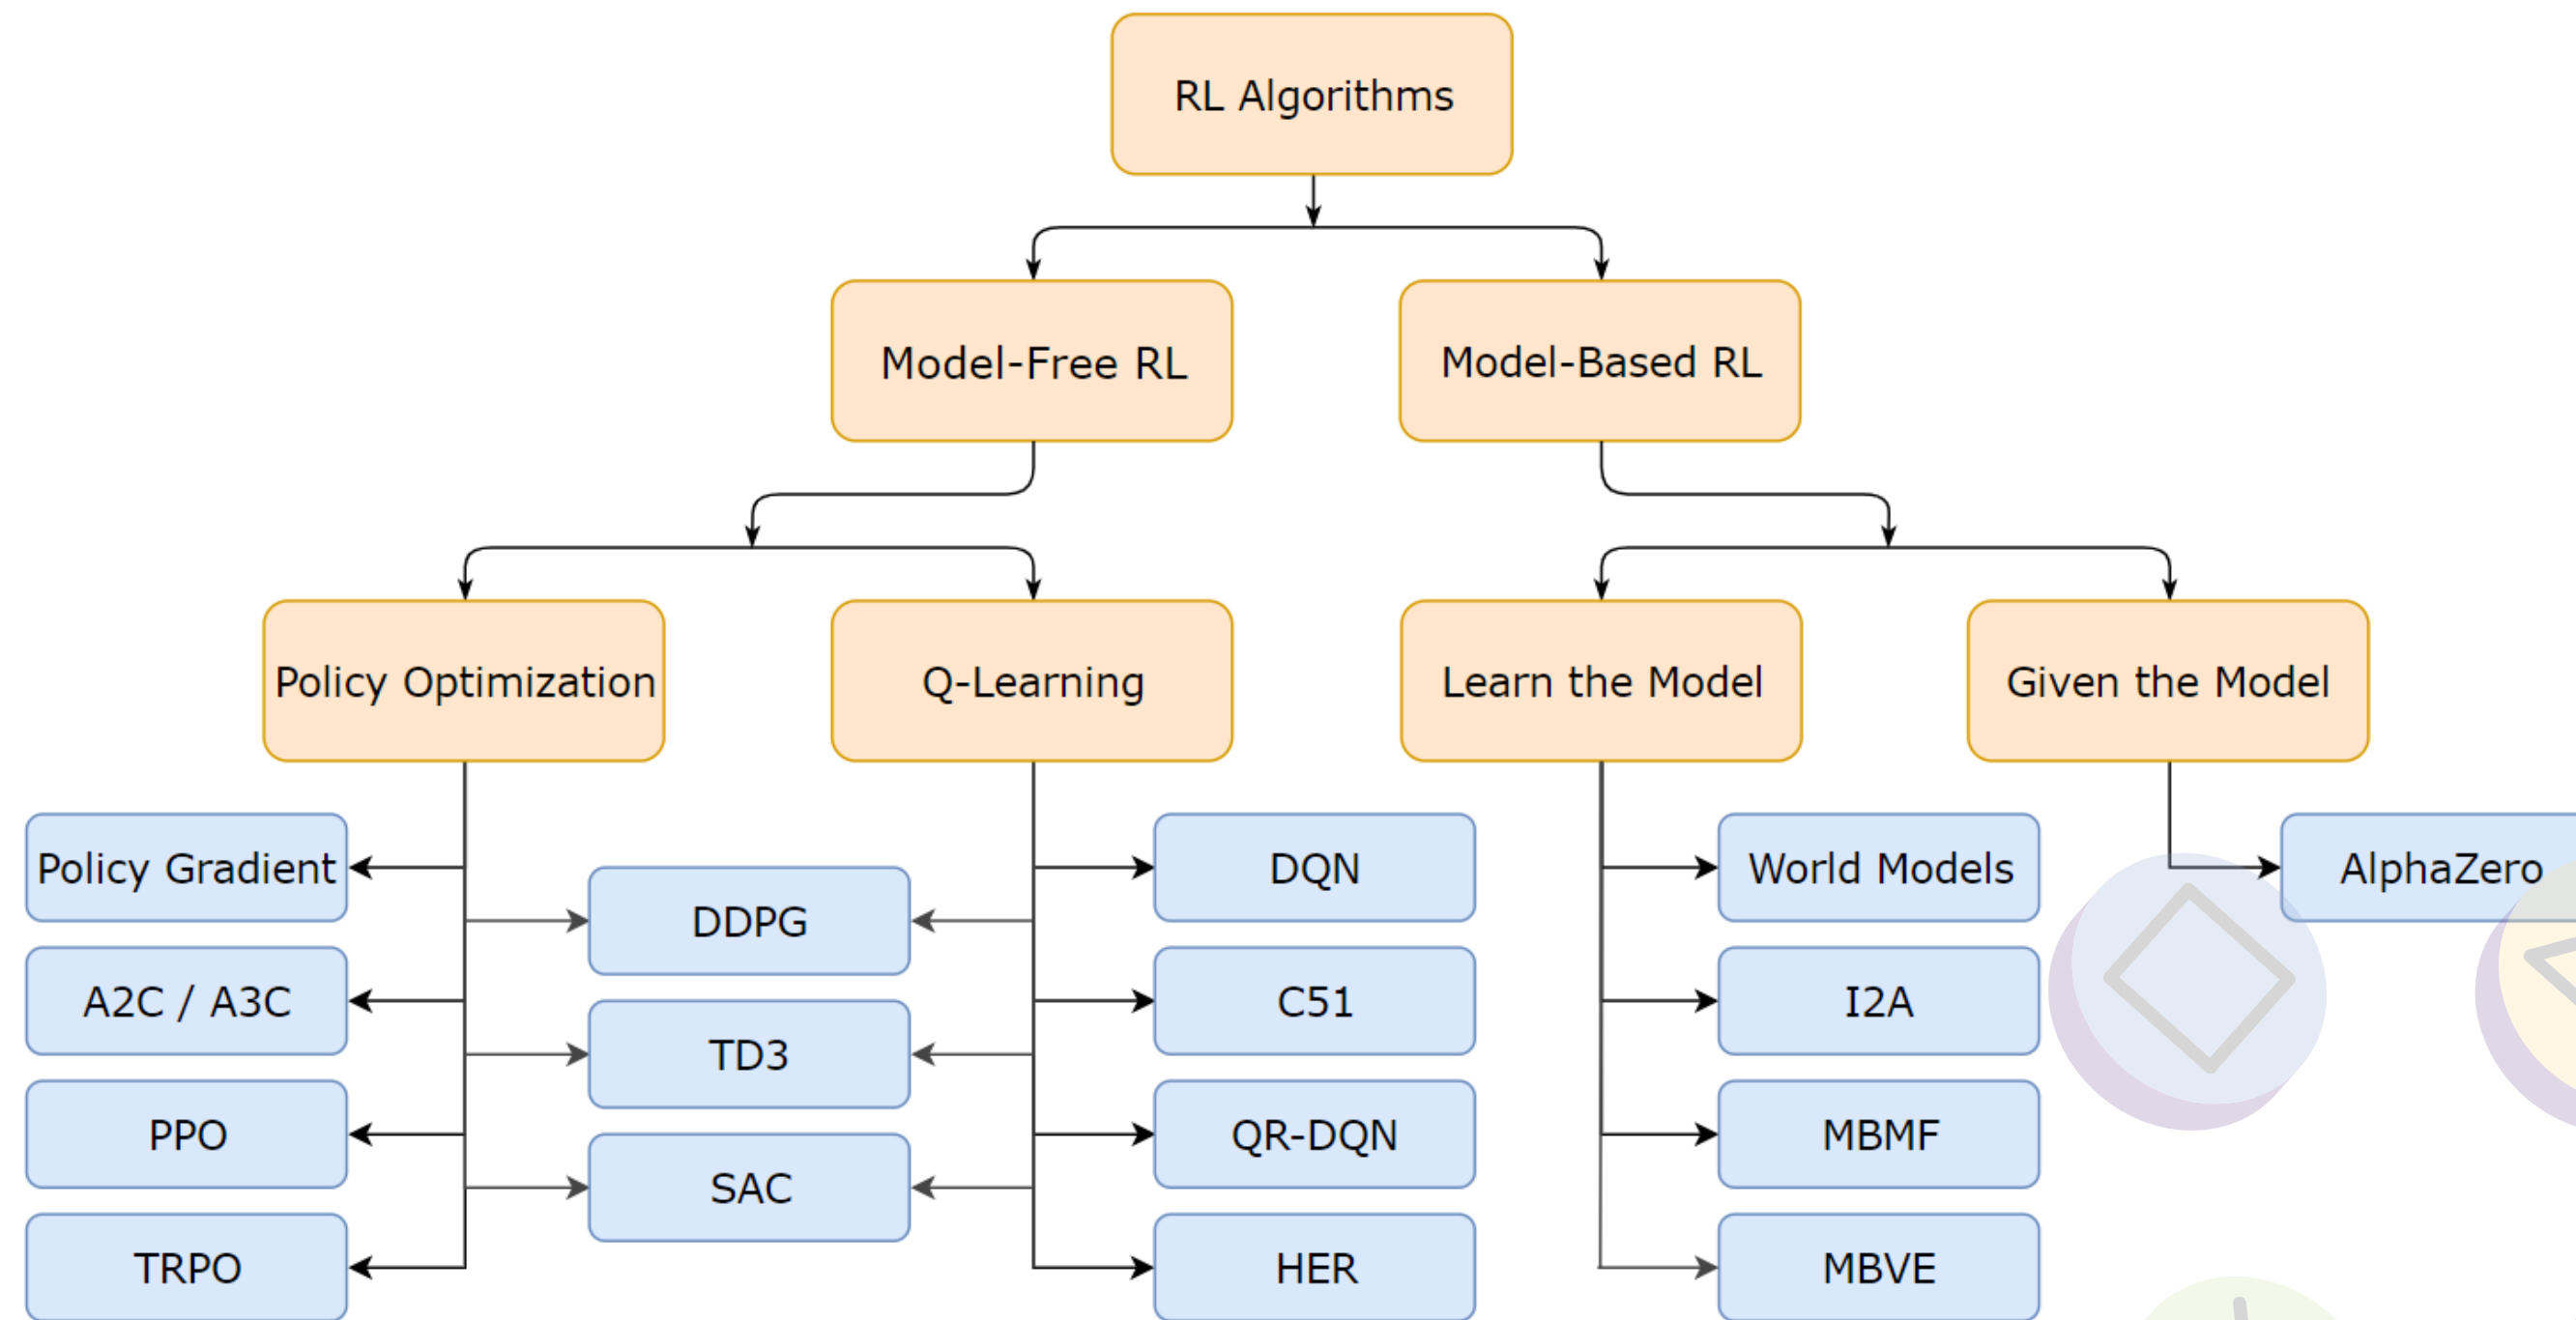


MODEL-FREE VS MODEL-BASED







- Model-Based: Agent can think ahead, and explicitly choose option by evaluating the possible choices.
- Model-Free: Easier to implement and tune.
 - Policy Optimization
 - Q-Learning



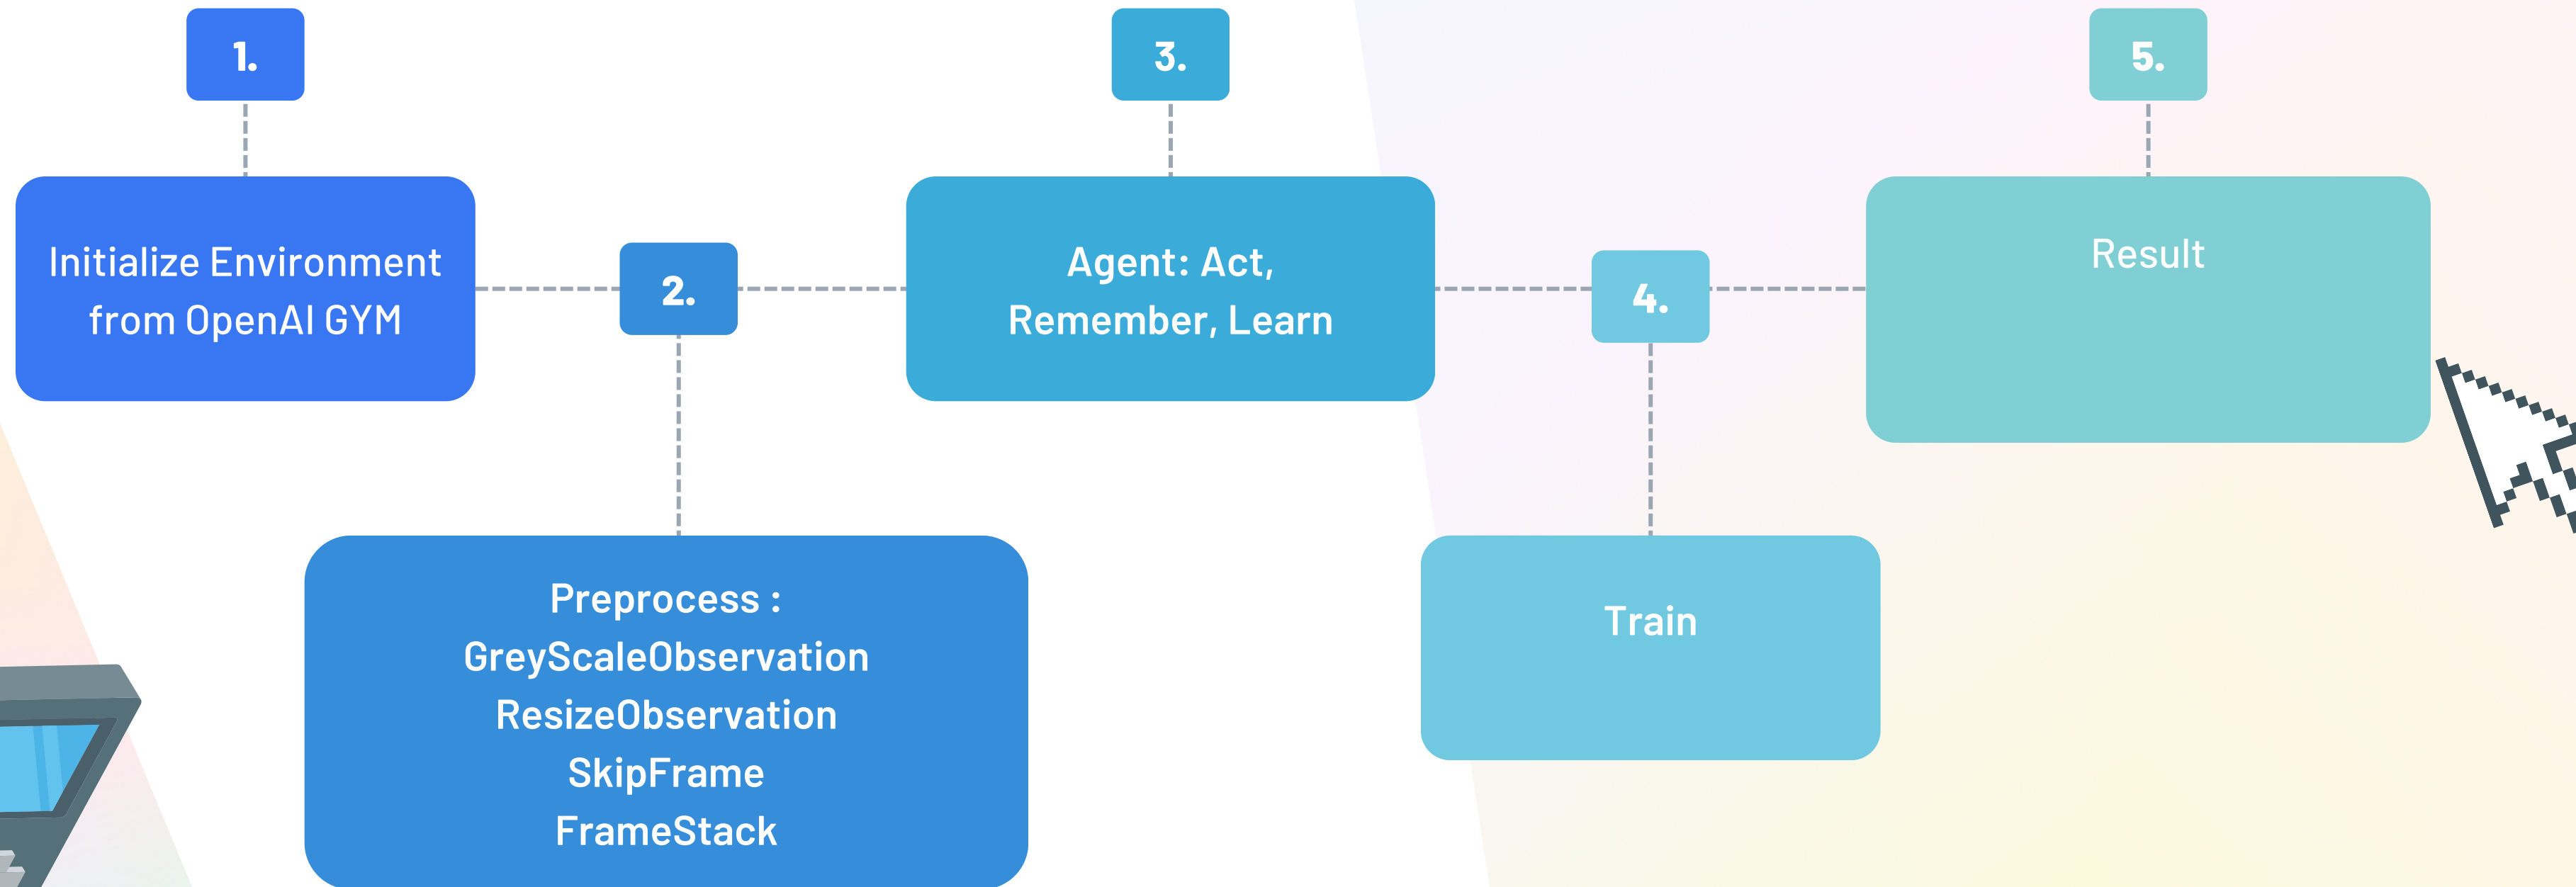
TYPES OF REINFORCEMENT LEARNING ALGORITHMS



MACHINE LEARNING TOOLS USED



EXPERIMENT SETUP



POLICIES


- Agent's Brain
- Deterministic or Stochastic
- Goal: Maximize Reward



$$a_t = \mu_{\theta}(s_t)$$

$$a_t \sim \pi_{\theta}(\cdot | s_t).$$

Parameterized Equation of Policies








POLICY GRADIENT

- Optimizes the parameters of policy directly
- Gradient based approach
- Two main steps: Policy Evaluation, Policy Improvement
- Example application domain: Control, Navigation, Robotics.





$$a_t = \mu_{\theta}(s_t)$$


$$a_t \sim \pi_{\theta}(\cdot | s_t).$$



POLICY GRADIENT - RESULT






POLICY GRADIENT - DEMO


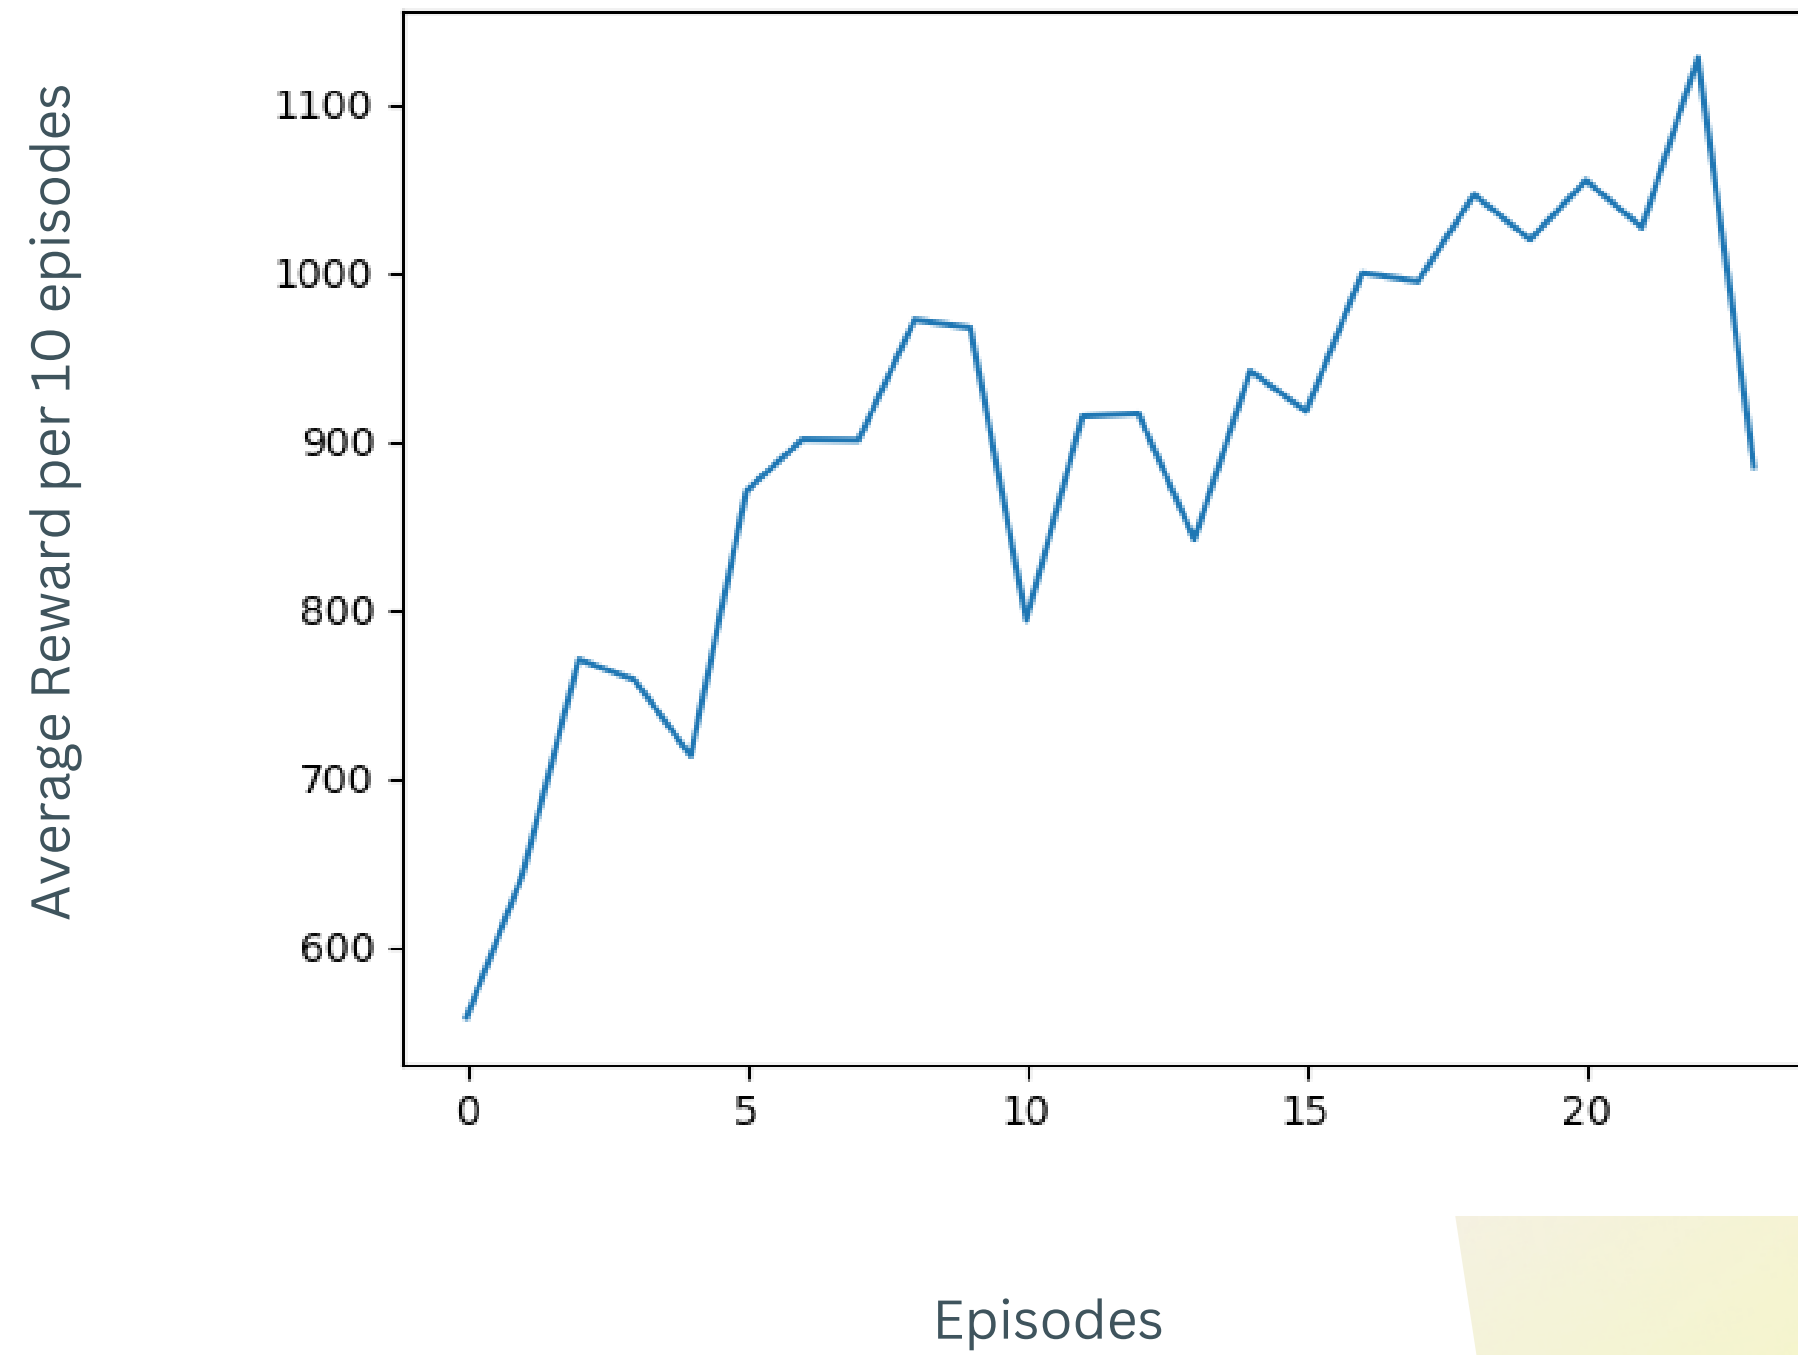


DEEP DOUBLE Q-LEARNING


- Combination of: Q-learning with Deep learning.
- Used two separate neural network.
- Replay buffer: Opportunity to learn from past.



DDQN - RESULT







ODQN - DEMO





PROXIMAL POLICY OPTIMIZATION (PPO)


- Aims to maximize policy performance.
- Clipped objective function: Prevent too much change at each iteration.
- Adaptive learning rate.
- Trust region constraint: Controls the step size of policy update.



PPO - RESULT





PPO - DEMO



LIMITATION & FUTURE WORK

- Loss plot, time plot.
- Other RL Algorithms.
- Experiment with learning rate, optimizer neural network layers.





THANK
YOU